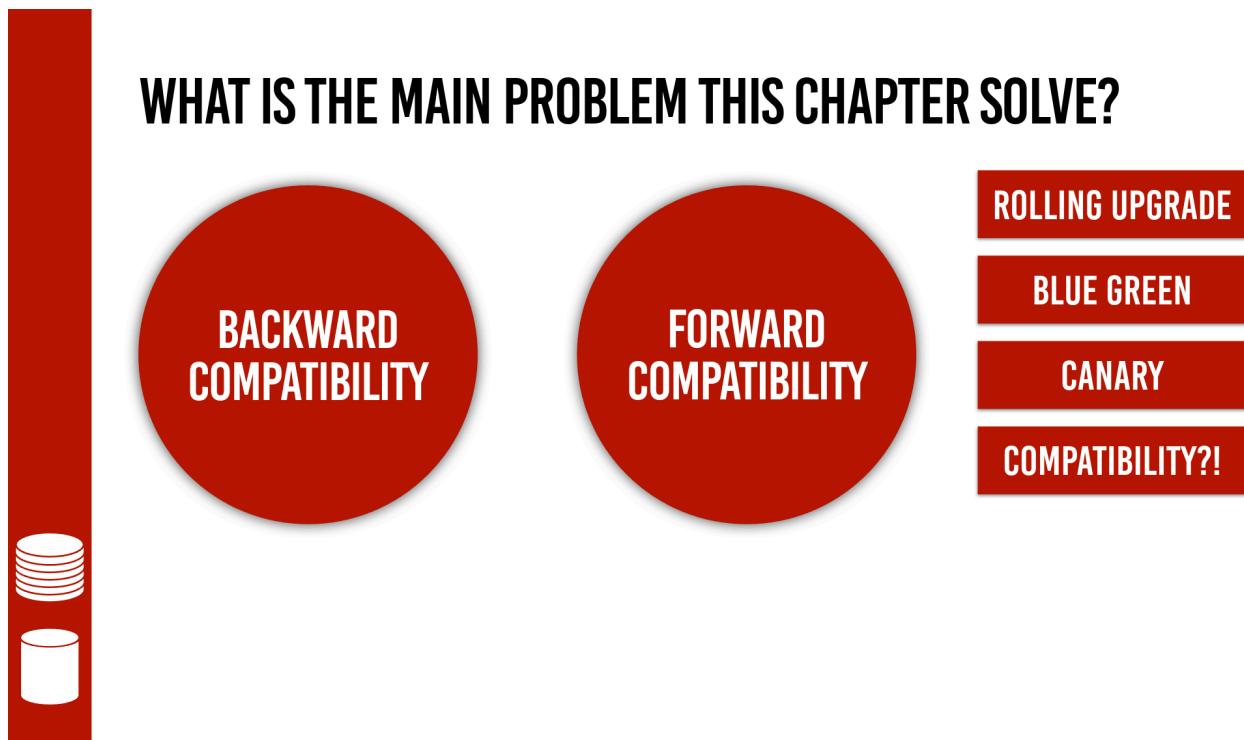


# What is the Main Problem this Chapter Solves?

- Explanation: The chapter addresses maintaining compatibility (backward: new code reads old data; forward: old code reads new data) during system evolution, using strategies like rolling upgrades, blue-green deployments, or canary releases to avoid downtime.
- Example: In social media platforms like Facebook, backward compatibility ensures old app versions can read new user data formats during gradual rollouts.
- Key Takeaway: - Compatibility ensures smooth updates in live systems, preventing breaks in dataflow across versions.



## Formats for Encoding Data

- Explanation: Lists common serialization formats (JSON, XML, Protocol Buffers, Thrift, Avro, Java Serializable) for converting in-memory data to bytes, solving the need for interoperable, efficient data transfer across processes.
- Example: In web services like REST APIs (e.g., Twitter), JSON encodes responses for browsers, while Avro handles internal Hadoop data for analytics.
- Key Takeaway: - Choosing the right format balances readability, compactness, and compatibility for reliable data-intensive apps.

# FORMATS FOR ENCODING DATA

SERIALIZATION

JSON

XML

PROTOCOL  
BUFFER

THRIFT

AVRO

JAVA.IO.SERIALIZABLE



## MessagePack Encoding Example

- Explanation: MessagePack is a binary format that converts structured data like JSON objects into a compact sequence of bytes for efficient storage or transmission, solving the problem of verbose text formats by packing data tightly without field names.
- Example: In Kafka, MessagePack can serialize event logs (e.g., user actions like clicks) into smaller payloads, reducing network bandwidth when streaming data between microservices.
- Key Takeaway: - MessagePack reduces data size for faster transfer in high-throughput systems like messaging queues.

```
{
  "userName": "Martin",
  "favoriteNumber": 1337,
  "interests": ["daydreaming", "hacking"]
}
```

MessagePack



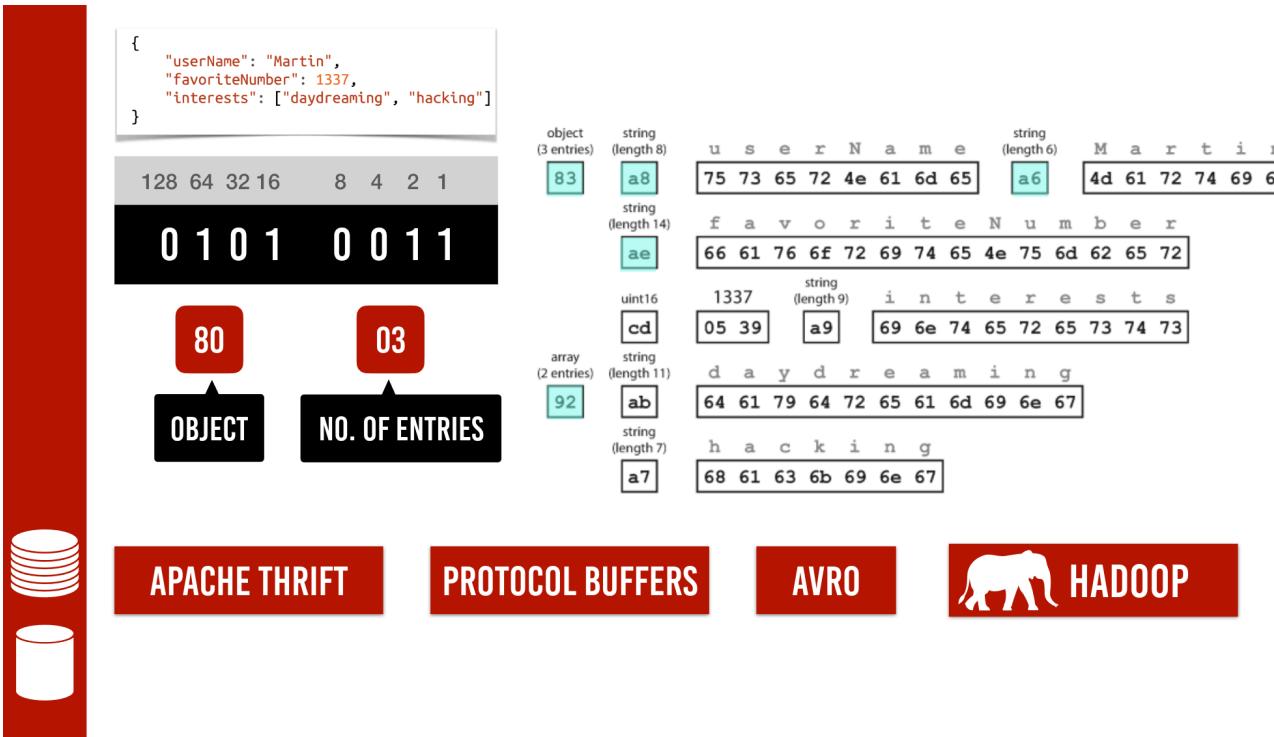
Byte sequence (66 bytes):

83	a8	75	73	65	72	4e	61	6d	65	a6	4d	61	72	74	69	6e	ae	66	61
76	6f	72	69	74	65	4e	75	6d	62	65	72	cd	05	39	a9	69	6e	74	65
72	65	73	74	73	92	ab	64	61	79	64	72	65	61	6d	69	6e	67	a7	68
61	63	6b	69	6e	67														



## Binary Encoding Comparison (Thrift, Protocol Buffers, Avro)

- Explanation: This compares how different binary formats (Thrift, Protocol Buffers, Avro) encode the same data object into bytes, showing varying sizes and structures; they use schemas to define data types, solving inefficiency in text formats like JSON by omitting redundant info.
- Example: In a social media app like Twitter, Avro encodes user profiles compactly for storage in Hadoop, while Protocol Buffers handles gRPC calls between services, saving space compared to JSON.
- Key Takeaway: - Binary formats with schemas enable compact, evolvable data encoding for large-scale systems like big data pipelines.



## Schemas for Thrift and Protocol Buffers

- Explanation: Schemas define data structures (e.g., required/optional fields) for Thrift and Protocol Buffers, solving ambiguity in data interpretation by providing a contract for encoding/decoding.
- Example: In a database like Cassandra integrated with Thrift, schemas define user record fields for queries, allowing evolution without rewriting all data.
- Key Takeaway: - Schemas facilitate evolvability, making systems adaptable to changes without data loss.



## APACHE THRIFT

```
struct Person {  
    1: required string      userName,  
    2: optional i64         favoriteNumber,  
    3: optional list<string> interests  
}
```

## PROTOCOL BUFFERS

```
message Person {  
    required string user_name      = 1;  
    optional int64  favorite_number = 2;  
    repeated string interests       = 3;  
}
```

BINARY PROTOCOL

COMPACT PROTOCOL

JSON PROTOCOL

COMPACT PROTOCOL



## Thrift BinaryProtocol Encoding

- Explanation: Thrift's BinaryProtocol encodes data with explicit type annotations and lengths for each field, using fixed tags instead of names, solving the verbosity of text while allowing schema evolution.
- Example: In microservices for e-commerce (e.g., Amazon), BinaryProtocol serializes order details for inter-service communication, ensuring compatibility during updates.
- Key Takeaway: - BinaryProtocol supports forward/backward compatibility, crucial for rolling upgrades in distributed systems.

## Thrift BinaryProtocol

Byte sequence (59 bytes): 

0b	00	01	00	00	00	06	4d	61	72	74	69	6e	0a	00	02	00	00	00	00
00	00	05	39	0f	00	03	0b	00	00	00	02	00	00	00	0b	64	61	79	64
72	65	61	6d	69	6e	67	00	00	00	07	68	61	63	6b	69	6e	67	00	

Breakdown:

type 11 (string) field tag = 1

0b      00 01

4 BYTES

length 6

M a r t i n

```
{  
    "userName": "Martin",  
    "favoriteNumber": 1337,  
    "interests": ["daydreaming", "hacking"]  
}
```

type 10 (i64) field tag = 2

0a      00 02

1337

00 00 00 00 00 00 05 39

8 BYTES

type 15 (list) field tag = 3

0f      00 03

item type 11 (string)

0b

2 list items

length 11

00 00 00 02

d a y d r e a m i n g

00 00 00 0b

64 61 79 64 72 65 61 6d 69 6e 67

length 7

00 00 00 07

h a c k i n g

end of struct

68 61 63 6b 69 6e 67

00

## Thrift CompactProtocol Encoding

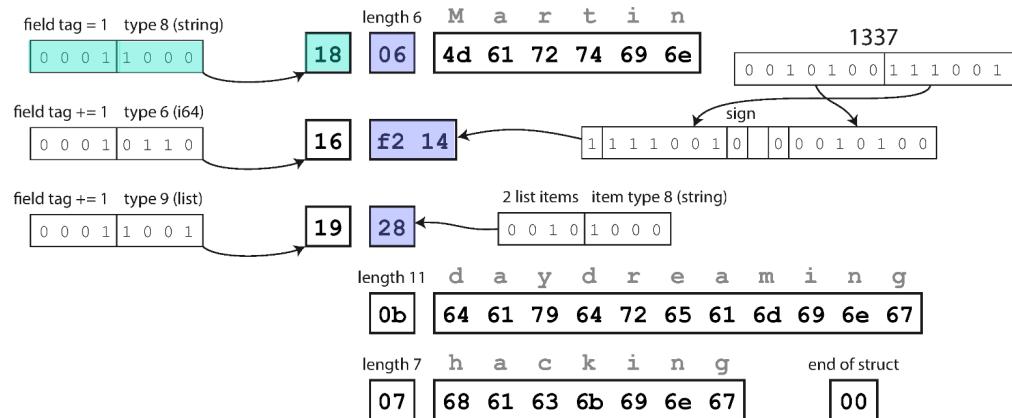
- Explanation: Thrift's CompactProtocol packs data into fewer bytes by combining field tags and types into single bytes and using variable-length integers, solving the space waste in simple binary encodings.
- Example: In a banking system, Thrift CompactProtocol serializes transaction records for RPC calls between fraud detection microservices, minimizing latency over the network.
- Key Takeaway: - CompactProtocol optimizes for space in bandwidth-constrained environments, like mobile-to-server communication.

## Thrift CompactProtocol

Byte sequence (34 bytes): 

18	06	4d	61	72	74	69	6e	16	f2	14	19	28	0b	64	61	79	64	72	65
61	6d	69	6e	67	07	68	61	63	6b	69	6e	67	00						

Breakdown:



## Protocol Buffers Encoding

- Explanation: Protocol Buffers (Protobuf) uses a schema with field tags to encode data compactly via bit-packing and variable integers, solving data evolution by allowing new fields without breaking old code.
- Example: In Google's internal systems or gRPC-based apps like ride-sharing (Uber), Protobuf encodes user location data for real-time updates, handling schema changes seamlessly.
- Key Takeaway: - Protobuf enables efficient, version-compatible data exchange in scalable services like APIs.

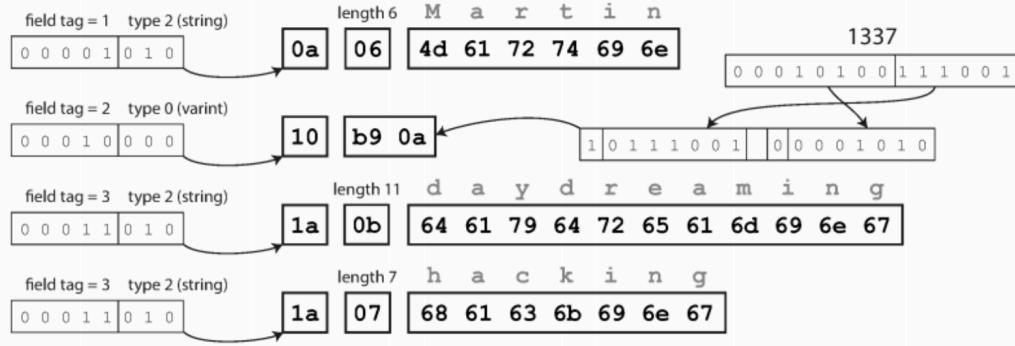


## Protocol Buffers

Byte sequence (33 bytes):

0a	06	4d	61	72	74	69	6e	10	b9	0a	1a	0b	64	61	79	64	72	65	61
6d	69	6e	67	1a	07	68	61	63	6b	69	6e	67							

Breakdown:



## Avro Encoding

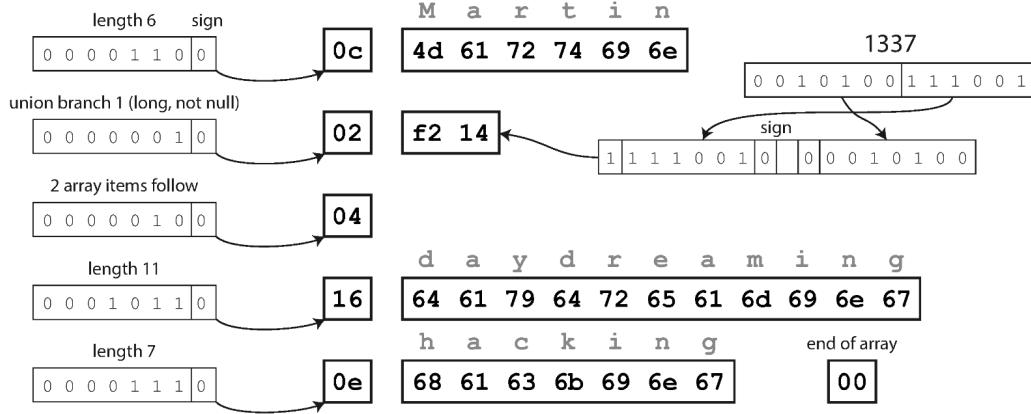
- Explanation: Avro encodes data using a schema resolved at read time (writer's vs. reader's schema), solving schema evolution by matching fields by name, not tags, and supporting defaults for missing data.
- Example: In Kafka streams for log processing (e.g., Netflix), Avro serializes events with evolving schemas, allowing consumers to handle old/new formats without disruption.
- Key Takeaway: - Avro's schema resolution supports dynamic changes in big data systems, ensuring long-term compatibility.

## Avro

Byte sequence (32 bytes):

0c	4d	61	72	74	69	6e	02	f2	14	04	16	64	61	79	64	72	65	61	6d
69	6e	67	0e	68	61	63	6b	69	6e	67	00								

Breakdown:



## WHAT IS AVRO & HOW IT WORKS?

- Explanation: Avro is a compact binary data format that uses a **schema** to describe the structure of data (like field names and types). When writing data, it follows the writer's schema and produces very small byte sequences. When reading, it uses the reader's schema and automatically matches fields by **name** (not by position or tag), filling in defaults for missing fields. This solves the problem of data breaking when schemas change over time.
- Example: Netflix uses Avro to serialize streaming events in Kafka. When they add a new field like "watchDurationMs" to the event schema, old consumers still read the data correctly (using defaults), and new consumers get the new field — no downtime or data loss.
- Key Takeaway:
  - Avro makes schema changes safe and automatic → ideal for long-lived data pipelines where producers and consumers evolve independently.

# WHAT IS AVRO & HOW IT WORKS?

EMAM

```
record Person {
    string          userName;
    union { null, long } favoriteNumber = null;
    array<string>   interests;
}

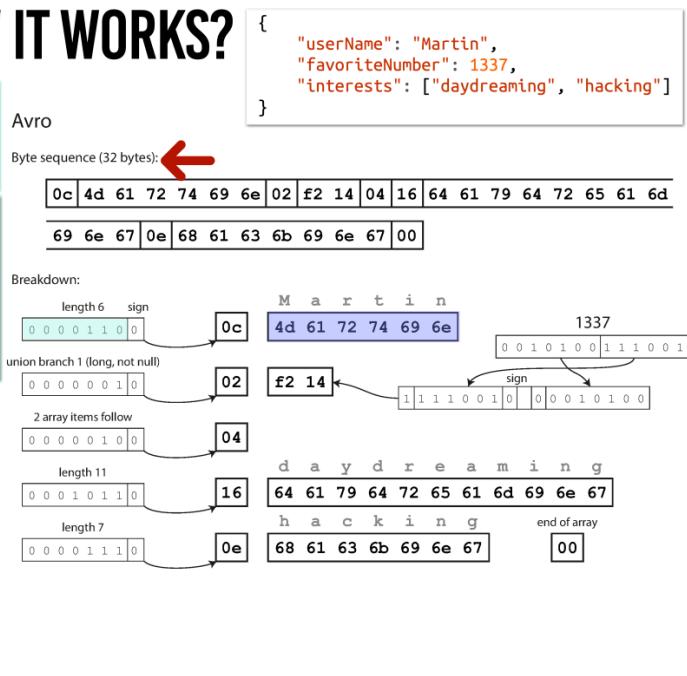
{
    "type": "record",
    "name": "Person",
    "fields": [
        {"name": "userName", "type": "string"},
        {"name": "favoriteNumber", "type": ["null", "long"]},
        {"default": null},
        {"name": "interests", "type": {"type": "array", "items": "string"}}
    ]
}
```

Thrift CompactProtocol

```
Byte sequence (34 bytes):
18 06 4d 61 72 74 69 6e 16 f2 14 19 28 0b 64 61 79 64 72 65
61 6d 69 6e 67 07 68 61 63 6b 69 6e 67 00
```

Breakdown:

field tag == 1 type 8 (string)	length 6	M a r t i n
0 0 0 0 1 0	0c	4d 61 72 74 69 6e
field tag == 1 type 0 (64)	1337	0 0 1 0 1 0 0 1 1 1 0 0 1
16	f2 14	union branch 1 (long, not null)
0 0 0 0 0 0 1 0	02	1 1 1 0 0 1 0   0 0 0 1 0 1 0 0
2 array items follow	04	sign
0 0 0 0 0 1 0 0	16	day d r e a m i n g
length 11	64 61 79 64 72 65 61 6d 69 6e 67	h a c k i n g
0 0 0 1 0 1 1 0	0e	length 7
length 7	68 61 63 6b 69 6e 67	end of array
0 0 0 0 1 1 1 0	00	

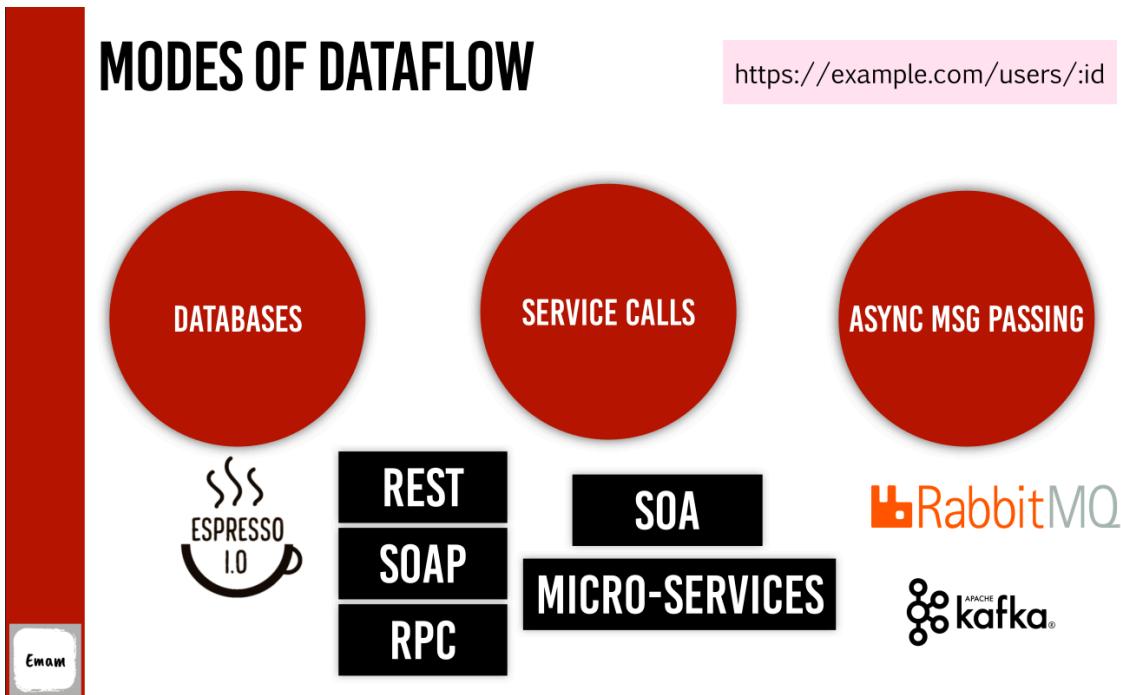


## MODES OF DATAFLOW

- Explanation: Data moves between systems in three main ways:
  - Databases** → write once, read later (possibly years later)
  - Service calls** (REST, SOAP, gRPC, RPC) → request → immediate response
  - Async message passing** (queues/topics like Kafka, RabbitMQ) → send message, someone consumes it later, no direct reply Each mode has different needs for encoding, compatibility, and performance.
- Example: In an e-commerce platform like Amazon:
  - Order is stored in a database (DynamoDB / Aurora)
  - Checkout service calls payment service via REST/gRPC
  - Order event is sent asynchronously via Kafka to warehouse, recommendation, and email systems
- Key Takeaway:
  - Understanding the three dataflow modes helps you choose the right serialization format and compatibility strategy for each part of a distributed system.

# MODES OF DATAFLOW

<https://example.com/users/:id>



## WHAT IS SCHEMA EVOLUTION?

- Explanation: Schema evolution means safely changing the data structure (adding/removing/rename fields, changing types) over time while old and new versions of code can still read/write the data correctly. Avro does this very well by matching fields **by name** and using defaults/unions → old code ignores new fields, new code fills defaults for missing old fields.
- Example: In LinkedIn's Espresso (their document store), they evolve user profile schemas frequently (add new fields like "pronouns", change types). Avro allows rolling upgrades: old servers keep writing/reading old format, new servers handle both — no full data migration needed.
- Key Takeaway:
  - Good schema evolution = you can deploy new application versions gradually without breaking old ones or losing data → critical for high-availability systems.

# WHAT IS SCHEMA EVOLUTION?

Writer's schema for Person record

Datatype	Field name
string	userName
union {null, long}	favoriteNumber
array<string>	interests
string	photoURL

Reader's schema for Person record

Datatype	Field name
long	userID
union {null, int}	favoriteNumber
string	userName
array<string>	interests

Εμμη