

WHAT IS DATA REPLICATION AND WHY?

- **Explanation:** Replication = keeping the same data on multiple machines. Three main reasons:
 - Reduce **latency** (serve users from nearby replica)
 - Increase **availability** (survive when one machine dies)
 - Increase **read throughput** (many replicas can answer read queries)
- **Example:** Netflix / YouTube → videos & metadata replicated in many regions → low latency + survive AWS zone failure.
- **Key Takeaway:** • Replication solves latency, availability, and read-scaling — the three biggest reasons to run more than one database node.

WHAT IS DATA REPLICATION AND WHY?



REDUCE
LATENCY

INCREASE
AVAILABILITY

INCREASE READ
THROUGHPUT

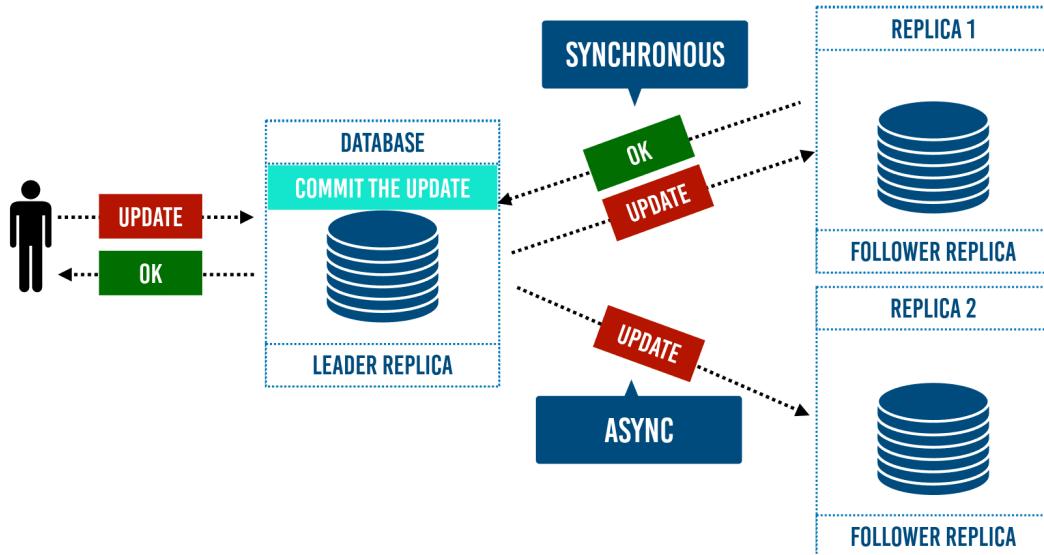
EMAIL

LEADERER-BASED REPLICATION S CHAPTER

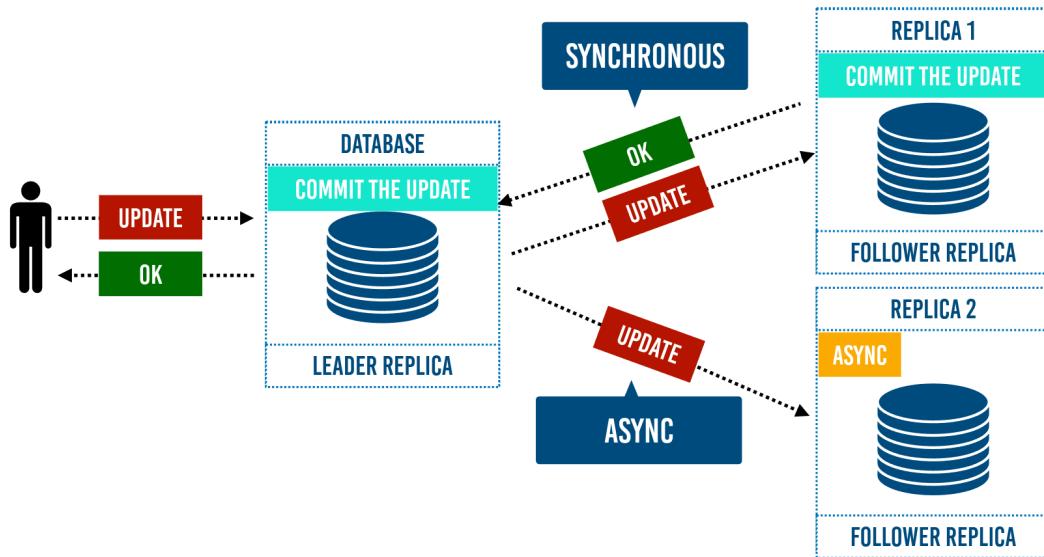
- **Explanation:** When the leader accepts a write:
 - **Synchronous** follower → leader waits until the follower confirms it received & applied the write before telling the client "OK".
 - **Asynchronous** follower → leader sends the change but does **not** wait — it tells the client "OK" immediately. → Synchronous = stronger durability, but slower & can block writes if follower is slow/down. → Asynchronous = faster, but you can lose recent writes if leader crashes.
- **Example:** PostgreSQL default = asynchronous replication (fast). You can turn on synchronous_commit = on → becomes semi-synchronous (leader waits for at least one synchronous standby → popular for important data).

- **Key Takeaway:** • Almost all production systems use **mostly asynchronous + 1 synchronous standby** (semi-synchronous) — best balance between speed and not losing data when leader dies.

LEADER-BASED REPLICATIONS CHAPTER?



LEADER-BASED REPLICATIONS CHAPTER?

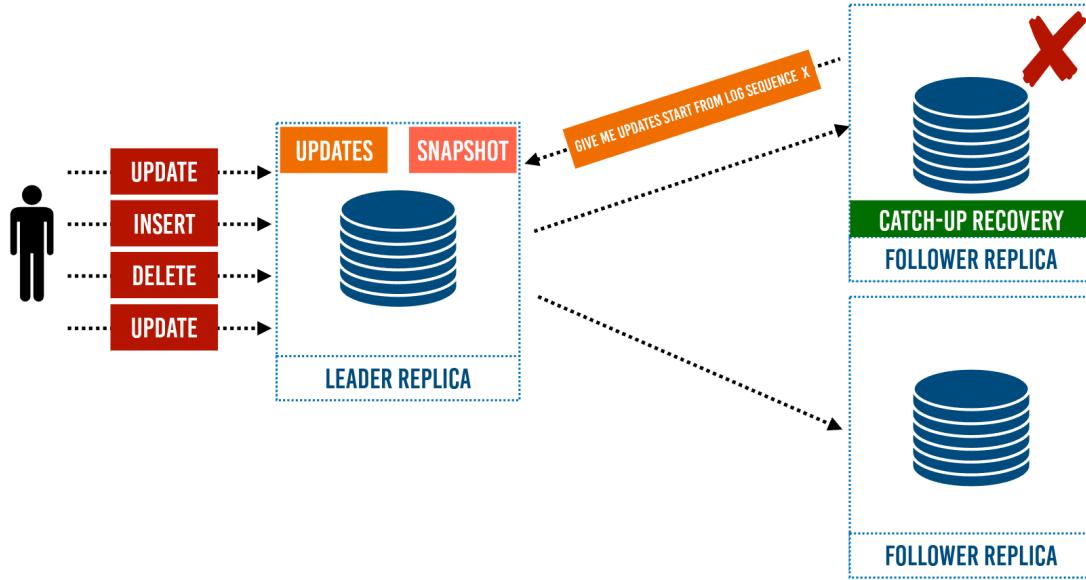


HOW TO BUILD NEW FOLLOWER REPLICA

- **Explanation:** Steps to add a new follower without downtime:
 1. Take consistent snapshot of leader (without locking everything).
 2. Copy snapshot to new follower.

3. Follower asks leader: "give me all changes since this snapshot position".
 4. Leader sends the backlog → follower applies it → now caught up and stays in sync.
- **Example:** PostgreSQL pg_basebackup + streaming replication. MySQL → Percona XtraBackup + binlog coordinates.
 - **Key Takeaway:** • Modern databases have automated / semi-automated ways to add new replicas without stopping writes — snapshot + catch-up is the standard pattern.

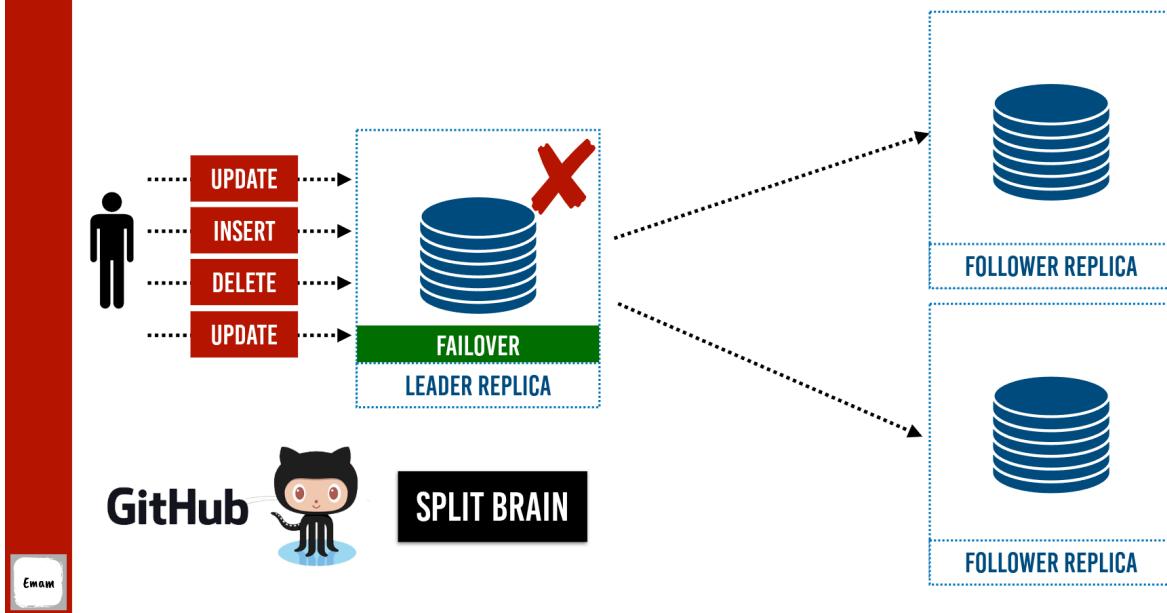
HOW TO BUILD NEW FOLLOWER REPLICA?



HOW TO BUILD NEW LEADER REPLICA?

- **Explanation:** When the current leader (primary) database dies or is unreachable, the system must promote one of the followers to become the new leader. But there's a dangerous risk: if the old leader comes back online later and still thinks it's the leader, both nodes might accept writes → this creates **split-brain** (two leaders thinking they are in charge → data corruption).
- **Example:** GitHub in 2012 had a famous split-brain incident during MySQL failover. An old primary came back online, both nodes accepted writes → same auto-increment IDs were reused → private user data leaked between accounts.
- **Key Takeaway:** • Split-brain is one of the most dangerous failure modes in leader-based replication — preventing it is why automatic failover is very hard and many teams prefer manual failover.

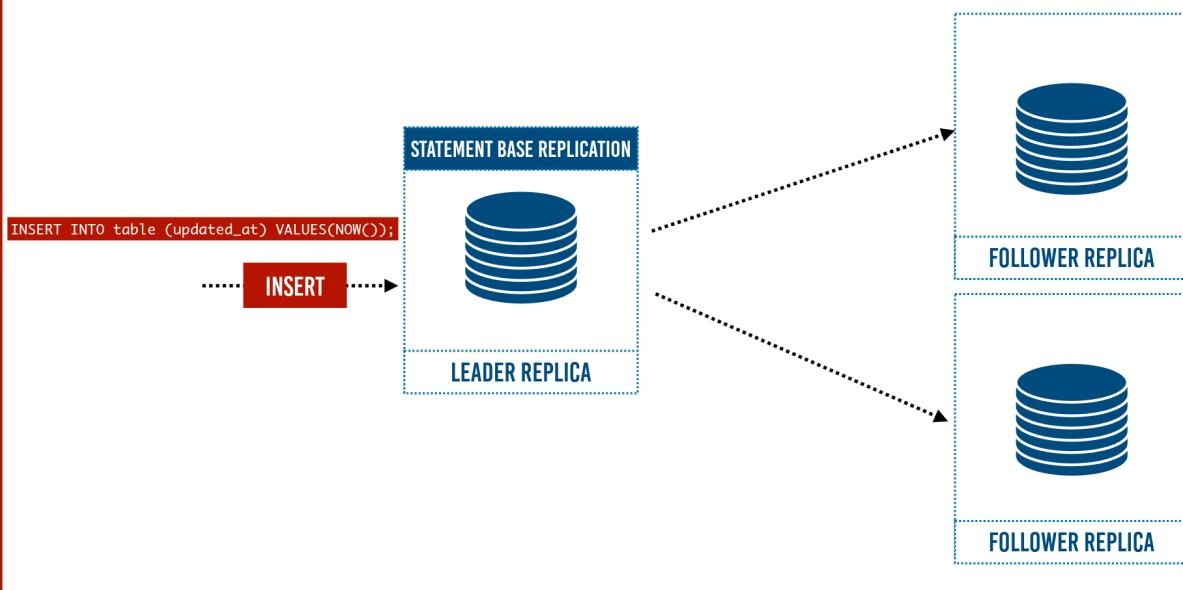
HOW TO BUILD NEW LEADER REPLICA?



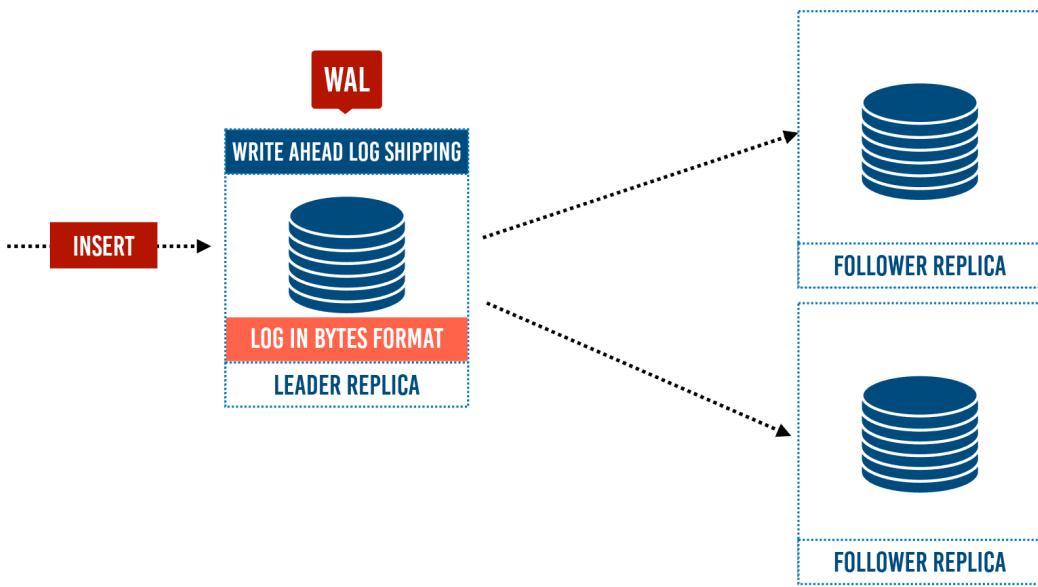
HOW DO LEADER SEND THE UPDATES TO FOLLOWERS

- **Explanation:** The leader doesn't send SQL statements or rows — it sends the **physical Write-Ahead Log (WAL)** or binlog entries (sequence of low-level changes: which bytes changed in which blocks). Followers receive the log and apply exactly the same changes → very efficient and keeps replicas identical.
- **Example:** PostgreSQL → WAL shipping (physical replication). MySQL → binlog (can be statement, row, or mixed format).
- **Key Takeaway:** • WAL / binlog shipping is the most common, battle-tested way leaders send changes — it's fast, reliable, and used by almost every production relational database.

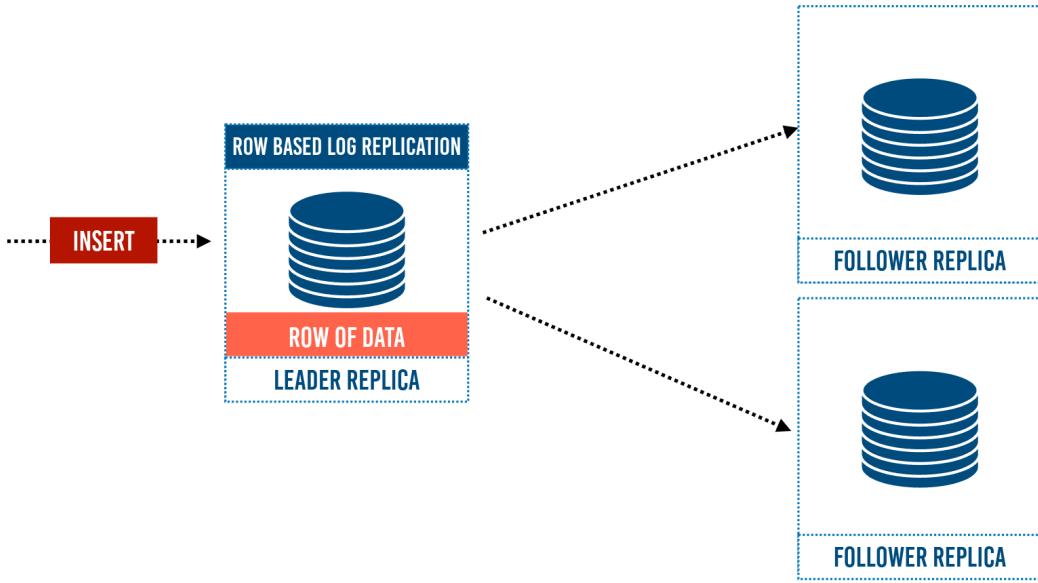
HOW LEADER SEND THE UPDATES TO FOLLOWERS?



HOW LEADER SEND THE UPDATES TO FOLLOWERS?



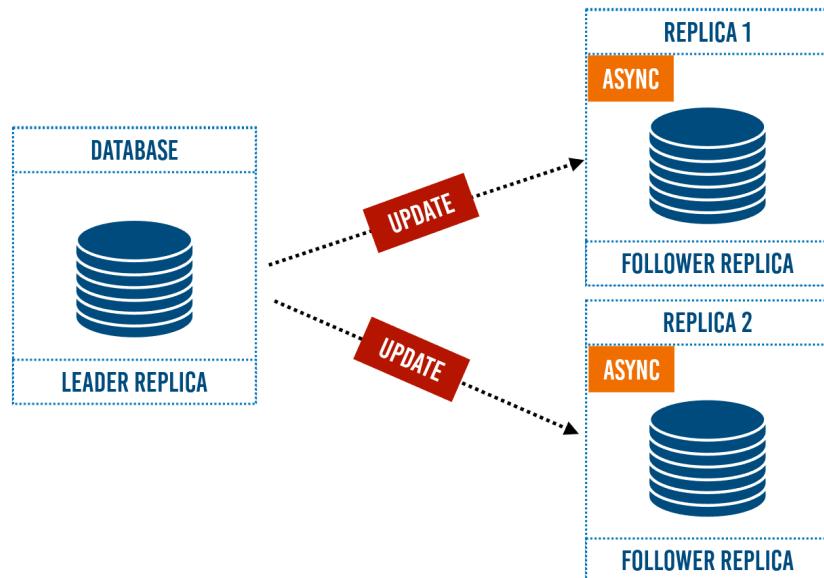
HOW LEADER SEND THE UPDATES TO FOLLOWERS?



LEADERER-BASED REPLICATION

- **Explanation:** Classic picture of leader-based replication: All writes → leader. Leader sends updates to many followers (usually async). Reads can go to leader or followers → read scaling.
- **Example:** Most web companies (Twitter, Instagram, Shopify...) use leader + many async followers for read-heavy traffic.
- **Key Takeaway:** • Leader-based + async replication = most common production pattern for relational databases today.

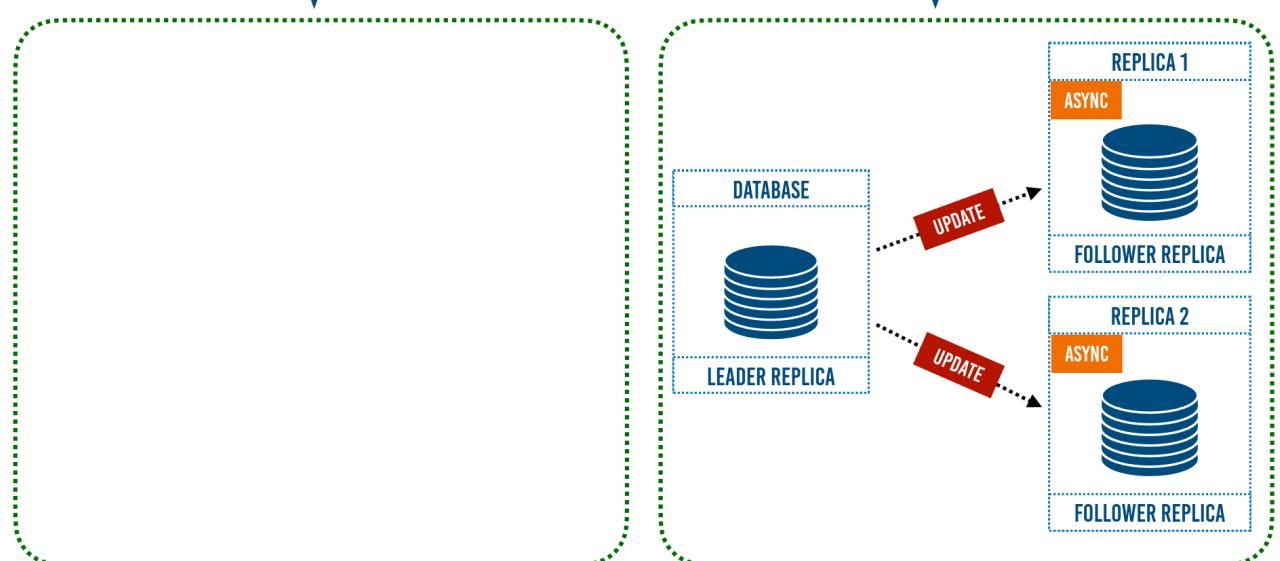
LEADER-BASED REPLICATION



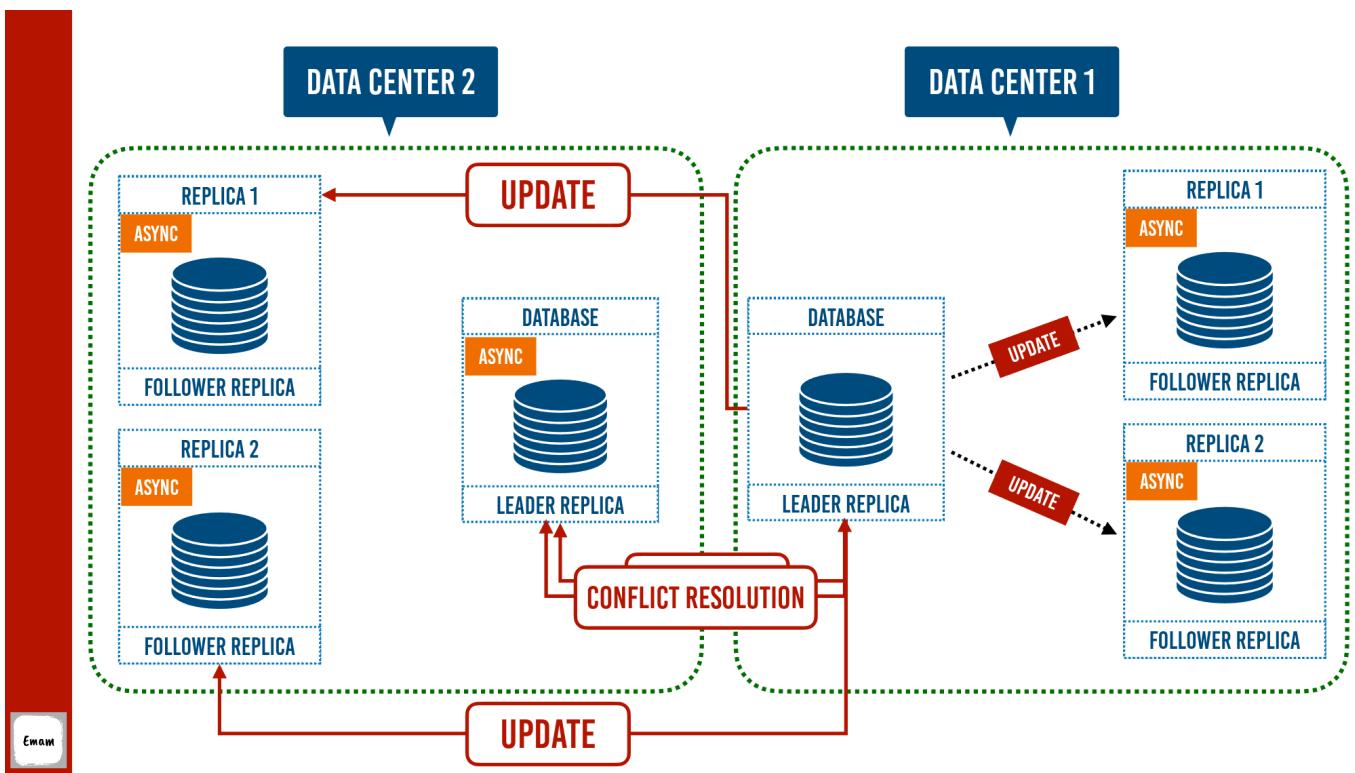
Émam

DATA CENTER 2

DATA CENTER 1



Émam



CONVERGENT CONFLICT RESOLUTION

- **Explanation:** Convergent = all replicas eventually reach the same value. Simplest (but dangerous) way: **Last Write Wins (LWW)** • Give every write a timestamp / unique ID • Keep only the one with highest number/ID → discard others
- **Example:** Cassandra default conflict resolution = LWW (very common). Many mobile apps also use LWW for simplicity.
- **Key Takeaway:** • LWW is easy but loses data silently — only use it when you can accept data loss (caches, non-critical counters...).

I covered the main distinct concepts/slides. If you want deeper explanation on any particular slide (e.g. OT algorithm steps, version vectors math, split-brain prevention techniques...), just

tell me which one.

CONVERGENT CONFLICT RESOLUTION

LAST WRITE
WIN

ON WRITE

ON READ

LWW

ÉMAM

CONVERGENT CONFLICT RESOLUTION

A/B

LAST WRITE
WIN

LWW

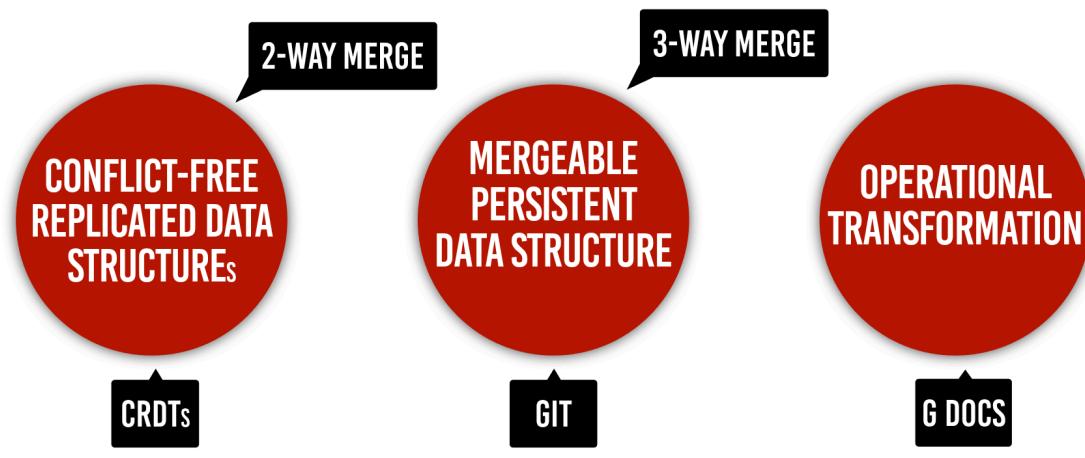
ÉMAM

AUTOMATIC CONFLICT RESOLUTION

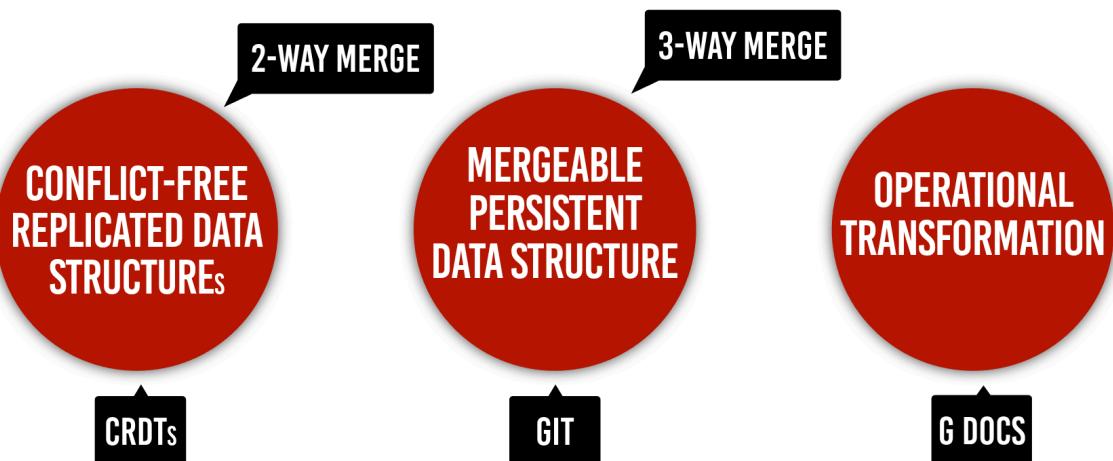
- **Explanation:** When the same data is changed in multiple places at once → conflict.
Automatic resolution approaches:
 - **CRDTs** — special data types that always merge correctly (counters, sets, lists...)
 - **Operational Transformation** (Google Docs style)
 - **3-way merge** (like Git)
 - **2-way merge** (simpler)

- **Example:** Riak → has CRDT counters/sets/maps. Git → 3-way merge for code conflicts. Google Docs → OT for text.
- **Key Takeaway:** • If you want multi-leader or offline-first without manual conflict handling → CRDTs or OT are the modern, clean answers.

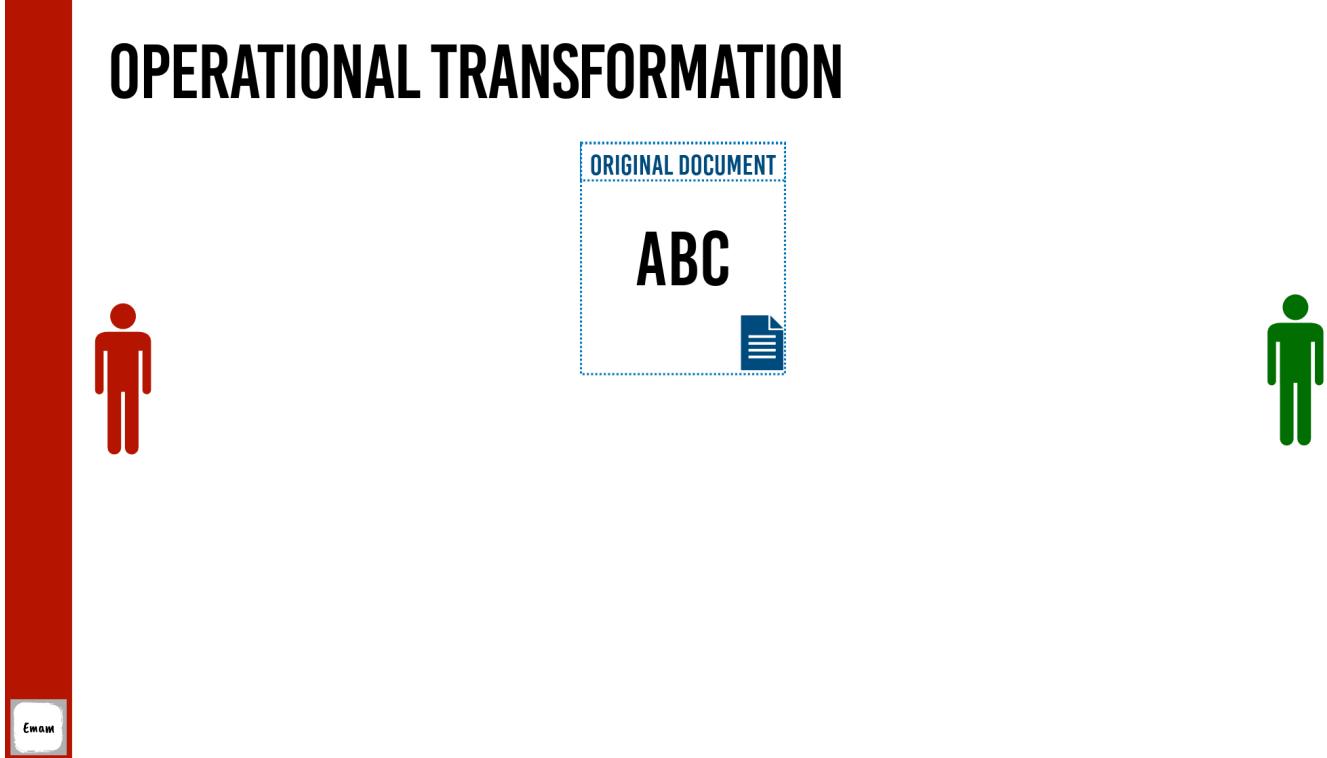
AUTOMATIC CONFLICT RESOLUTION



AUTOMATIC CONFLICT RESOLUTION



OPERATIONAL TRANSFORMATION



OPERATIONAL TRANSFORMATION

- **Explanation:** Operational Transformation (OT) is a technique used when multiple people edit the **same document at the same time** (Google Docs, Etherpad...). Instead of sending whole document, you send small operations (insert char X at position 5, delete char at 7...). When operations arrive in different order on different clients → OT algorithm **transforms** them so everyone ends up with the same final text (even though operations were applied in different sequence).
- **Example:** Google Docs: you and your colleague type at the same time → your cursor doesn't jump around wildly — OT transforms your colleague's insert so it appears correctly relative to what you just typed.
- **Key Takeaway:** • OT is the magic behind real-time collaborative editing — it turns concurrent operations into consistent document state without locking.

OPERATIONAL TRANSFORMATION



Email



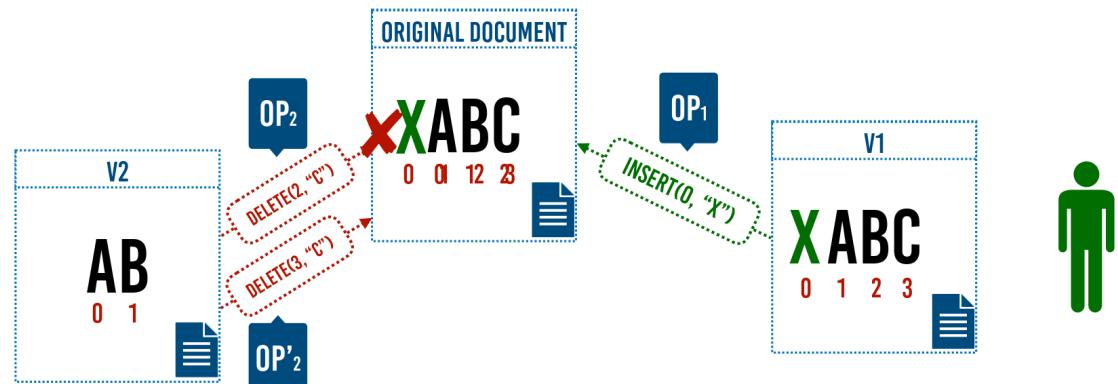
OPERATIONAL TRANSFORMATION



Email



OPERATIONAL TRANSFORMATION



EMAIL